Retail Inventory Optimization Agent

Decision-Grade AI System for Weekly Replenishment & Working-Capital Prioritization

Author: Taash Chikosi

Target roles: AI Consultant / AI Product Manager / AI Solutions Architect

Target employers: Top-tier consulting firms and large enterprises adopting AI at scale

Value proposition: Design and deliver end-to-end AI agent systems that solve real business problems, quantify ROI, and integrate human-in-the-loop decision-making

Strengths: Business problem framing, AI workflow design, ML-driven predictions, RAG, executive dashboards

Outcome focus: Measurable efficiency gains, cost reduction, or revenue impact

Table of Contents

1. Executive Summary

This engagement delivers a live **AI decision-support system** designed to help retail buyers and planners **prioritize replenishment actions and deploy limited working capital more effectively** during weekly purchasing cycles.

The system forecasts near-term SKU-store demand based on current conditions and recent sales patterns, then translates those forecasts into **standard inventory policies**—safety stock, reorder points, and recommended order quantities. These outputs are further reframed into **economic trade-off signals** that explicitly surface capital exposure, holding-cost risk, and stockout-avoidance benefits.

The system is intentionally designed for **decision support and prioritization**, not for automated purchasing or financial reporting. Its purpose is to help decision-makers answer a single commercially critical question early enough to matter:

“Given current demand and supply conditions, where should we deploy working capital this week to reduce stockout risk—and which recommendations require human review?”

The deployed solution combines:

- a machine-learning demand forecasting layer trained on realistic retail-like operational data, and
- a retrieval-augmented reasoning (RAG) layer that explains outputs, assumptions, limitations, and governance boundaries using a curated retail knowledge base.

2. Decision Context & Objectives

Decision context

Retail inventory decisions are made under persistent uncertainty. Demand varies by product, store, season, and promotion; supplier lead times fluctuate; and working capital is always constrained.

While retailers collect extensive historical sales and inventory data, decision-makers often lack a **forward-looking, decision-grade signal** that:

- reflects the current operational state,

- translates forecasts into concrete buying actions, and
- makes capital trade-offs explicit.

The core challenge is not perfect demand prediction, but **better prioritization of attention and capital**.

Business question

Based on current sales trends, inventory position, and supplier lead times:

Which SKUs require replenishment this week, how much should be ordered, and which actions justify human review due to capital or risk exposure?

Intended users

- Retail buyers — order planning and approval
- Replenishment and demand planners — prioritization and exception management
- Operations teams — service-level and availability monitoring
- Finance partners — working-capital visibility and trade-off assessment

Success criteria

- Forward-looking demand signals tied to current conditions
- Translation of forecasts into familiar inventory policy constructs
- Explicit economic framing of capital vs service-level trade-offs
- Conservative, interpretable outputs suitable for weekly decision-making
- Clear boundaries on appropriate and inappropriate use

3. How the System Works (High Level)

The system follows a deliberately simple and defensible workflow:

Observe → Forecast → Translate → Prioritize → Explain → Decide

Inputs (current operational state only)

Users do not provide future assumptions. The system operates on observable inputs available at decision time:

- Recent daily sales history (SKU–store)
- Current on-hand inventory
- Outstanding purchase orders
- Product attributes (category, cost, margin proxies)
- Supplier lead-time statistics
- Calendar context (day-of-week and seasonal effects)

Recent rolling demand statistics are computed conservatively to reflect realistic historical context without requiring full enterprise data feeds.

Outputs

For each SKU–store requiring attention, the system produces:

- Predicted near-term demand
- Safety stock requirement
- Reorder point
- Recommended order quantity
- Working capital required (proxy)
- Holding cost exposure (proxy)
- Stockout penalty avoided (proxy)
- Net value prioritization signal
- Evidence-backed explanations and governance guidance via the RAG assistant

Outputs are framed to **support buyer judgment**, not automate purchasing.

4. What the Results Show

For a representative retail scenario, the system produces:

- A ranked list of SKUs requiring replenishment
- Clear differentiation between low-impact and high-impact actions
- Explicit visibility into capital exposure for each recommendation
- A conservative estimate of economic trade-offs associated with acting or not acting

Importantly, results are **directional rather than deterministic**, answering:

“Where should we focus first this week?”

The system does not prescribe budgets, override buyers, or optimize profit automatically.

5. Why the Results Are Trustworthy

This system was designed to be **decision-grade**, not merely predictive. Trustworthiness is established through disciplined data design, conservative modeling, and explicit governance controls.

Data strategy and realism

Access to real enterprise retail data is often restricted. The dataset was therefore engineered to **behave like real retail operations**, not idealized simulations.

Key principles included:

- Separation of demand, inventory, and supply dynamics
- Inclusion of demand volatility and noise
- Realistic lead-time variability
- Avoidance of steady-state or best-case bias

Synthetic data generation and ML readiness

Synthetic data was generated through a reproducible pipeline that preserved causal relationships (e.g., demand → inventory depletion → replenishment decisions) while introducing uncertainty and measurement error.

The dataset was designed to be:

- suitable for supervised learning,
- representative of retail decision contexts, and
- aligned with time-aware validation.

Feature engineering and unit of analysis

The unit of analysis is **one SKU at one store at one point in time**, forecasting near-term demand and decision implications.

All features are derived from:

- observable current conditions, or
- recent historical summaries available at decision time.

This ensures:

- no look-ahead bias,
- deployment realism, and
- alignment with actual buying workflows.

Training, validation, and performance criteria

Models were trained using a temporal split, ensuring training data precedes test data in time.

Performance evaluation prioritized:

- Mean Absolute Error (MAE) as a practical measure of forecast deviation,
- stability across SKUs and stores, and
- robustness under noisy inputs.

A tree-based regression model was selected for its balance of performance, stability, and interpretability in operational settings.

Governance, confidence, and human oversight

Governance controls are explicit and conservative:

- Clear separation between forecasts and decisions
- Rule-based translation from prediction to action
- Mandatory human review under high-capital or high-uncertainty scenarios
- Language constraints to prevent over-claiming
- Logged inputs and outputs for auditability

Together, these mechanisms ensure the system **supports human judgment rather than replacing it**.

6. What Is Live & What Comes Next

What is live today

- Deployed Stream lit web application with a Retail decision tab
- Version-controlled ML inference artifacts
- Persisted RAG vector store with retail-specific knowledge
- Secure secrets management and reproducible deployment

Users can interactively explore replenishment decisions and observe how capital and risk trade-offs change.

What the deployment proves

- End-to-end operability
- Stable, low-latency inference
- Buyer-friendly outputs for non-technical users
- Practical integration of governance and explainability

Scope boundaries

The system is designed for:

- weekly replenishment planning
- decision support and prioritization
- exception-based escalation

The system is not designed for:

- automated purchasing
- real-time inventory control
- financial reporting or accounting use
- regulatory or compliance certification

What comes next

Potential enhancements include:

- scenario comparison across service levels and capital constraints
- promotion-aware demand regimes
- override logging for learning loops
- category-level aggregation views
- integration with live enterprise data feeds

Closing Summary

This Retail Inventory Optimization Agent demonstrates how AI can be applied responsibly to improve **commercial decision-making** in complex retail environments.

By combining realistic data design, disciplined forecasting, explicit economic translation, governance controls, and evidence-backed explanation, the system delivers **defensible, decision-grade signals** that help retail leaders act with clarity—without overstating certainty or automating judgment.